

Prior-Authorization Automation

OCR + agent + human-in-the-loop across payer portals with no clean API

The challenge

Clinic staff spend hours a day on payer portals, faxes and prior-auth packets. Manual, slow, and error-prone — and a denied or delayed auth means delayed patient care.

Generic RPA breaks on OTP, bot-detection and layout changes.

-70%

staff hours on auth work

72h → 6h

auth turnaround

+22%

first-pass approvals

Auth turnaround time

Before



After



Solution architecture

Intake: fax / PDF / portal

OCR + field extraction

Eligibility rules (code)

Agent fills portal + human approve

Financial-Document Extraction with an Accuracy Gate

Structured, trustworthy data — with the eval that proves it before it ships

The challenge

Lending ops key bank statements and tax docs by hand.
Errors are expensive and slow underwriting — but nobody trusts automation without proof the numbers are right.

So accuracy has to be measured, not assumed.

99.2%

field accuracy (from 91%)

-60%

cost per document

80%

docs auto-cleared

Field-level extraction accuracy

Before



91%

After



99.2%

Solution architecture

Docs in

OCR + schema-enforced extract

Confidence gate

Accuracy / regression gate → LOS

Support Agent that Deflects — Safely

Grounded, cited answers + guardrails + human sign-off on anything that moves money

The challenge

Support volume spikes at peak. Off-the-shelf bots invent policy, leak customer data, and escalate the wrong tickets — so the team turns them off and drowns in the queue.

Deflection only counts if the answers are safe and correct.

48%

tickets deflected

4.6/5

CSAT on bot chats

0

policy hallucinations

Tickets resolved without an agent

Before

0%

After

48%

Solution architecture

Customer query

RAG over policies (cited)

Scoped order/refund tools

Guardrails + human approval

Contract Review with Source Citations

Citation-grounded retrieval — no citation, no answer

The challenge

Associates spend hours hunting clauses across long contracts.
A naive AI assistant answers confidently and is sometimes wrong — which in legal work is worse than no answer at all.

Every answer must be traceable to the source clause.

-65%

clause-review time

100%

answers cite a source

0

uncited answers shown

Time to review a contract

Before



100 min

After



35 min

Solution architecture

Contracts


Chunk + provenance

Retrieval w/ score floor

Cited answers + gold-set eval


Citation-Grounded RAG

Retrieval that refuses to guess — the fix for confident-but-wrong document assistants




Reliable Multi-Agent Harness

One codebase, N agents — adding an agent is configuration, not new code




Secure MCP Server

Letting an AI agent drive your write APIs — safely (scopes, tenant isolation, idempotency)




LLM Evaluation & Regression Gating

Make model quality measurable — and stop silent quality drift before it ships



AI Automation / Document Extraction

Unstructured input to trustworthy structured output — with humans on the risky steps



LLM Cost & Performance Optimization

Profile first, fix the real bottleneck — production-ready and scalable, without a rewrite

